

---

---

# OpenMP Dialect

FC + MLIR

---

---

# Overview

- Dialect to represent OpenMP constructs
- FC AST nodes are converted to OpenMP dialect operations.
- `omp.parallel`, `omp.single`, `omp.do`, `omp.master` and `omp.parallel_do` are currently supported
- Data elements to be mapped are represented as arguments to the operations.
- Nesting of constructs are supported for parallel do, etc.
- Parallel regions are outlined to functions at LLVM Dialect level

# Implementation Details

- AST Nodes Definition : **include/AST/StmtOpenMP.h**
- Parser support : **lib/parse/ParseOpenMP.cpp**
- MLIR code generation: **lib/codegen/CGOpenMP.cpp**
- OpenMP dialect operations: **include/dialect/OpenMP/OpenMPOps.td**
- LLVM Dialect lowering: **lib/transforms/OpenMPLowering.cpp**
- OpenMP unit tests: **test/openmp/**

# OpenMP Dialect Operations

(include/dialect/OpenMP/)

# Parallel construct

```
a = 10
```

```
b = 20
```

```
!$omp parallel
```

```
print *, "hellofrom", a, b
```

```
!$omp end parallel
```

```
omp.parallel(%2, %3)  {  
    %5 = fc.constant_string("hellofrom")  
    %6 = fc.load %2 {name = "a"} : i32  
    %7 = fc.load %3 {name = "b"} : i32  
    fc.print %5, %6, %7  
}
```

# Parallel Do Construct

```
!$omp parallel do
```

```
do i = 1, k
```

```
    c(i) = a(i) + b(i)
```

```
end do
```

```
!$omp end parallel do
```

```
omp.parallel_do(%5 = %6, %7, %8) (%4, %2, %3) {  
    %9 = fc.load %5 {name = "i"} : i32  
    %10 = index_cast %9 : i32 to index  
    %11 = fc.load %2[%10] {name = "a", range_info =  
        #fc.subscript_range< 0 >} : i32  
    %12 = fc.load %5 {name = "i"} : i32  
    %13 = index_cast %12 : i32 to index  
    %14 = fc.load %3[%13] {name = "b", range_info =  
        #fc.subscript_range< 0 >} : i32  
    %15 = addi %11, %14 : i32  
    %16 = fc.load %5 {name = "i"} : i32  
    %17 = index_cast %16 : i32 to index  
    fc.store %15, %4[%17] {name = "c", range_info =  
        #fc.subscript_range< 0 >} : !fc.ref<!fc.array<1:10 x  
        i32>>  
} enddo
```

# Do construct

- omp do construct is represented using OMP::OmpDoOp.
- Upper bound, lower bound and step are represented as arguments to OmpDoOp operation.

```
!$omp do
```

```
do i = 1, 10
```

```
    print *, "Hello omp do "
```

```
enddo
```

```
!$omp end do
```

```
omp.do %arg0 = %c1_i32, %c10_i32, %c1_i32
```

```
    {construct_name = ""} {
```

```
        %2 = fc.constant_string("Hello omp do ")
```

```
        fc.print %2 {arg_info = #fc.is_string< 0 >}
```

```
    } enddo
```

# Single construct

```
integer  :: n
n = 10

!$omp single
  print *, "Hello from omp", n
!$omp end single
```

```
%c10_i32 = constant 10 : i32

omp.single() {
  %3 = fc.constant_string("Hello from omp")
  fc.print %3, %c10_i32 {arg_info = #fc.is_string< 0 >}
}
```



# Master construct

```
!$omp master
```

```
print *, "Hello from master"
```

```
!$omp end master
```

```
print *, "From outside"
```

```
omp.master()  {  
    %3 = fc.constant_string("Hello from master")  
    fc.print %3 {arg_info = #fc.is_string< 0 >}  
}  
  
%2 = fc.constant_string("From outside")  
fc.print %2 {arg_info = #fc.is_string< 0 >}
```

# Nesting of constructs (1) : Parallel do

```
!$omp parallel do
```

```
do i = 1, k
```

```
    !$omp parallel do
```

```
do j = 1, k
```

```
    c(i, j) = a(i, j) + b(i, j)
```

```
end do
```

```
    !$omp end parallel do
```

```
end do
```

```
!$omp end parallel do
```

```
omp.parallel_do(%5 = %11, %12, %13) (%6, %7, %4, %2, %3) {
```

```
    ...
```

```
    omp.parallel_do(%6 = %15, %16, %17) (%7, %4, %5, %2, %3) {
```

```
        ...
```

```
        fc.store %28, %4[%30, %32] {name = "c"}
```

```
    } enddo
```

```
} endo
```

# Nesting of constructs(2)

```
!$omp parallel
```

```
    print *, "Hello from omp"
```

```
    $omp single
```

```
        print *, "Hello from single thread"
```

```
    !$omp end single
```

```
!$omp end parallel
```

```
omp.parallel() {
```

```
    %3 = fc.constant_string("Hello from omp"):fc.array<0:14 x i8>
```

```
    fc.print %3 {arg_info = #fc.is_string< 0 >}
```

```
omp.single() {
```

```
    %4 = fc.constant_string("Hello from single thread")
```

```
    fc.print %4 {arg_info = #fc.is_string< 0 >}
```

```
}
```

```
}
```

# Matmul Example

```
!$omp parallel do
```

```
do i = 1,n
```

```
do j = 1,n
```

```
c(i,j) = 0
```

```
do k = 1,n
```

```
c(i,j) = c(i,j) + a(i,k) * b(k,j)
```

```
enddo
```

```
enddo
```

```
enddo
```

```
!$omp end parallel do
```

```
omp.parallel_do(%3 = %9, %10, %11) (%4, %2, %5, %0, %1) {
```

```
...
```

```
fc.do %arg0 = %c1_i32_9, %c300_i32_10, %c1_i32_11 {
```

```
...
```

```
fc.do %arg1 = %c1_i32_13, %c300_i32_14, %c1_i32_15 {
```

```
...
```

```
fc.store %26, %2[%28, %arg0] {name = "c"}
```

```
} enddo
```

```
} enddo
```

```
} enddo
```

# OpenMP Dialect Lowering

(lib/transforms/OpenMPLowering.cpp)

# Dialect Lowering

- OpenMP dialect lowering is scheduled after all the other dialects are lowered to LLVM.
- Operations are lowered directly to LLVM IR.
- Outlines OpenMP parallel regions to `LLVM::FuncOp`
- Outlined function is called using OpenMP runtime library function `__kmpc_fork_call(..)`

# ParallelOp Lowering

```
omp.parallel(%2) {  
  
  %2 = fc.load %2 {name = "a"} : i32  
  
  fc.print %2 {arg_info =  
    #fc.is_string< >}  
  
}
```

```
define void @outlined.(i32* noalias nocapture readnone %0, i32* noalias  
  nocapture readnone %1, i32* noalias nocapture readonly %2) {  
  %4 = load i32, i32* %2, align 4, !alias.scope !0, !noalias !2  
  tail call void (i32, ...) @__fc_runtime_print(i32 2, ,i32 3, i32 %4)  
  ret void  
}  
  
define i32 @foo() local_unnamed_addr {  
  %1 = alloca i32, align 4  
  store i32 10, i32* %1, align 4  
  call void (@ i32, i32, i32, i32, i8* }*, i32, i8*, ...)  
    @__kmpc_fork_call(@ i32, i32, i32, i32, i8* }* nonnull  
    @ident.global, i32 1, i8* bitcast (void (i32*, i32*, i32*)*  
    @outlined. to i8*), i32* nonnull %1)  
  
  ret i32 0  
}
```

# SingleOp Lowering

```
%c10_i32 = constant 10 : i32
```

```
omp.single() {
```

```
  %3 = fc.constant_string("Hello from  
omp")
```

```
  fc.print %3, %c10_i32
```

```
}
```

```
%1 = call i32 @__kmpc_global_thread_num({ i32, i32, i32, i32, i8* }*  
                                           nonnull @ident.global)
```

```
%2 = call i32 @__kmpc_single({ i32, i32, i32, i32, i8* }* nonnull  
                             @ident.global, i32 %1)
```

```
%3 = icmp eq i32 %2, 0
```

```
br i1 %3, label %5, label %4
```

```
4:                                     ; preds = %0
```

```
  call void (i32, ...) @__fc_runtime_print(i32 4, i32 9, i8* getelementptr  
10)  inbounds ([16 x i8], [16 x i8]* @str_const_2, i64 0, i64 0), i32 3, i32
```

```
  call void @__kmpc_end_single({ i32, i32, i32, i32, i8* }* nonnull  
                               @ident.global, i32  
%1)
```

```
  br label %5
```

```
5:                                     ; preds = %0, %4
```

```
  call void @__kmpc_barrier({ i32, i32, i32, i32, i8* }* nonnull  
                           @ident.global,  
i32 %1)
```





# OmpDoOp Lowering

```
omp.do %arg0 = %c1_i32, %c10_i32, %c1_i32 {  
    %2 = fc.constant_string("Hello omp do ")  
    fc.print %2  
} enddo
```

```
%5 = call i32 @__kmpc_global_thread_num({ i32, i32, i32, i32, i8* }*  
    nonnull @ident.global)  
call void ({ i32, i32, i32, i32, i8* }*, i32, i32, i32*, i32*, i32*, i32,  
    i32, ...) @__kmpc_for_static_init_4({ i32, i32, i32, i32, i8* }*  
    nonnull @ident.global, i32 %5, i32 34, i32* nonnull %4, i32*  
    nonnull %1, i32* nonnull %2, i32* nonnull %3, i32 1, i32 1)  
.....  
br i1 %9, label %._crit_edge, label %.lr.ph  
.lr.ph:                                ; preds = %0, %lr.ph  
.... ; Loop body  
._crit_edge:                          ; preds = %lr.ph, %0  
call void @__kmpc_for_static_fini({ i32, i32, i32, i32, i8* }* nonnull  
    @ident.global, i32 %5)  
call void @__kmpc_barrier({ i32, i32, i32, i32, i8* }* nonnull  
    @ident.global, i32 %
```



**Thank You**

